

Agent Chat Handling

Admin Tracking & Oversight Guide

Platform: Frankie (Agent Nebula frontend) + Furdeco-ChatBot (.NET backend)

Audience: Supervisors / Admins

Purpose: How an admin monitors live chat handling, enforces accountability, and measures performance — verified across the customer, agent and admin surfaces.

1 Tracking philosophy

Admin oversight is built on three layers:

- **Real-time visibility** — every live conversation, the queue, and who is / isn't working right now.
- **Accountability** — automatically surface agents who are free while customers wait, or who don't respond in time.
- **Historical measurement** — per-agent performance, customer ratings and SLA outcomes over time.

Everything below is wired end-to-end (customer widget → backend → agent app → admin console), so the figures reflect real activity, not mock data.

2 Real-time visibility (Admin → Live Chats)

2.1 Team Health Strip

A live row of tiles, refreshed via SignalR + a 15s poll:

Tile	Meaning	Highlights
Waiting	Customers currently in the queue	amber when > 0
Longest wait	Oldest waiting customer (minutes)	red past 5 min
SLA breaches	Queued customers past the threshold	red when > 0
Online	Agents marked available	—
In chats	Agents handling ≥ 1 chat	—
Idle	Online agents handling zero chats	red while customers wait

One glance = whole-team health. High “Waiting” + red “Idle” means capacity is being wasted.

2.2 Idle-while-waiting banner (key accountability signal)

When agents are **Online but handling no chats** while customers wait, a red banner appears **naming those agents**:

■ **2 agents free** while 5 customers waiting — Aditya Tripathi, Agent One

2.3 Live conversation oversight

- Read the transcript live (customer + agent, real time).
- Per-session timer (queued / accepted / status).
- Full **Order details** — consignment, recipient, address, delivery, status, items, tracking — captured at chat start.
- **Whisper** a private tip; **Barge in** to take over; **End chat** (creates a ticket + rating prompt).

3 Accountability & routing (auto-assign)

Configured in **Admin** → **Settings** → **Auto-assign**. Off by default (pull mode).

3.1 Pull mode (default)

Chats sit in the Session Queue; agents claim them via **Accept next** (FIFO) or **per-row Accept** (cherry-pick). The Health Strip exposes anyone idling while the queue grows.

3.2 Auto-assign mode (push)

- Queued chats are pushed to the **longest-idle** available agent (under the cap).
- Agent must send a **first reply within the response window** (default 30s).
- No reply & another agent available → **reassigned + agent flagged**; otherwise it stays put (no orphaning).
- After **max attempts** with everyone tried → **escalated** to admins for manual pickup.
- **One pending offer per agent** — chats are fed one at a time, never piled on one agent.

3.3 Settings (admin-editable at runtime)

Setting	Default	Range
Enable auto-assign	Off	on / off
Max chats / agent	5	1–20
Response window (s)	30	10–300
Max attempts before escalation	3	1–10

4 Alerts & notifications (Admin bell)

Pushed in real time, persisted, and de-duplicated (raised once):

Notification	Trigger
New Chat	A customer joins the queue
Queue-wait breach	A customer waits past the threshold (default 5 min)
SLA breach	A ticket passes its SLA deadline unresolved
Assign timeout	(auto-assign) Agent didn't respond; chat reassigned
Auto-assign escalation	No agent responded after max attempts
Chat transferred	A chat moved between agents

5 Historical performance (Admin → Analytics)

Per-agent leaderboard, driven by real session data:

Column	What it tells you
Handled / Resolved	Volume and completion
Resolution %	Closure rate
Avg first reply	Responsiveness (time-to-first-message)
Avg CSAT	Customer satisfaction (from star ratings)

Customer ratings (1–5★ after every chat) surface as Avg CSAT per agent / workspace, plus a star badge on each ticket — so a 1★ chat is easy to find and follow up. Every resolved chat becomes a **ticket** carrying the transcript, agent, rating, SLA status and order ref.

6 Implementation map (customer · agent · admin)

Customer (chatbot widget): single-composer live chat with emoji + End Chat; two-way typing; captures the order snapshot at chat start; rates the agent 1–5★; sees supervisor join + colour-coded supervisor messages on barge.

Agent (Live Chats app): status (persists across refresh) driving Idle/Online metrics; Inbox + Session Queue (Accept next / cherry-pick); receives auto-assigned chats (auto-opens, “respond now”); sees the customer's order details; activity feeds the leaderboard.

Admin (console): Live Chats (Health Strip, idle banner, transcripts, whisper, barge, end chat, order panel); Settings → Auto-assign; Analytics leaderboard + CSAT; Agents (status / active chats / last seen); Tickets; Notifications.

7 Recommended admin monitoring routine

- **Continuous** (Live Chats open): watch the Health Strip; if the idle banner turns red, nudge the named agents or enable auto-assign.
- **On alert**: act on Queue-wait / Escalation notifications immediately.
- **Daily**: review the Analytics leaderboard (low first-reply / CSAT / resolution %) and spot-check 1★ tickets.
- **Capacity tuning**: if the queue backs up, raise staffing, the per-agent cap, or enable auto-assign.

8 Implementation status checklist

#	Capability	Status	Where
1	Live transcript oversight (all active chats)	Done	Admin → Live Chats
2	Team Health Strip	Done	Admin → Live Chats
3	Idle-while-waiting banner (named)	Done	Admin → Live Chats
4	Whisper / Barge / End chat	Done	Admin → Live Chats
5	Order details panel (admin + agent)	Done	Live Chats / agent profile
6	Queue-wait alerts	Done	Background service → bell
7	SLA-breach alerts	Done	Background service → bell
8	Per-agent performance leaderboard	Done	Admin → Analytics
9	Customer CSAT (ratings → badges + averages)	Done	Tickets / Analytics
10	Auto-assign (round-robin, cap, timeout, escalate)	Done	Settings + engine
11	Agent status (persists across refresh)	Done	Agent app + profile
12	Session Queue (Accept next + cherry-pick)	Done	Agent + Session Queue

Deployment note: features relying on backend logic (auto-assign, queue-wait alerts, order snapshot on the agent endpoint, status persistence) require the backend redeploy; the UI pieces ship with the frontend build. Both are built and ready.

